

ML in Applications

***Dipartimento di Automatica e Informatica
Politecnico di Torino, Torino, ITALY***



**Politecnico
di Torino**

Lab 4

Tensorflow - Introduction



Tensors and variables

- TensorFlow operates on multidimensional arrays called **tensors** objects
- On tensors tf implements:
 - Canonical math operations
 - ML specialized ops
- Tensors are **immutable**
- To store model weights TensorFlow use a **tf.Variable**

```
[11] 1 import tensorflow as tf
      2
      3 x = tf.constant([[1., 2., 3.],
      4                  [4., 5., 6.]])
      5
      6 print(x)
      7 print(x.shape)
      8 print(x.dtype)
```

```
tf.Tensor(
[[1. 2. 3.]
 [4. 5. 6.]], shape=(2, 3), dtype=float32)
(2, 3)
<dtype: 'float32'>
```

```
[12] 1 x + x
```

```
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[ 2.,  4.,  6.],
       [ 8., 10., 12.]], dtype=float32)>
```

```
[13] 1 tf.nn.softmax(x, axis=-1)
```

```
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=
array([[0.09003057, 0.24472848, 0.66524094],
       [0.09003057, 0.24472848, 0.66524094]], dtype=float32)>
```

Automatic differentiation

- Optimization through **gradient descent** is a ML milestone
- To compute gradients tf implements automatic differentiation

```
[15] 1 x = tf.Variable(1.0)
      2
      3 def f(x):
      4     y = x**2 + 2*x - 5
      5     return y
      6
      7 f(x)
```

```
<tf.Tensor: shape=(), dtype=float32, numpy=-2.0>
```

```
[16] 1 with tf.GradientTape() as tape:
      2     y = f(x)
      3
      4 g_x = tape.gradient(y, x) # g(x) = dy/dx
      5
      6 g_x
```

```
<tf.Tensor: shape=(), dtype=float32, numpy=4.0>
```

The [tf.GradientTape](#) API computes the gradient with respect to some inputs

Automatic differentiation


- Gradient of a loss with respect to two variables

```
[3] 1 import tensorflow as tf
    2
    3 w = tf.Variable(tf.random.normal((3, 2)), name='w')
    4 b = tf.Variable(tf.zeros(2, dtype=tf.float32), name='b')
    5 x = [[1., 2., 3.]]
    6
    7 with tf.GradientTape(persistent=True) as tape:
    8     y = x @ w + b
    9     loss = tf.reduce_mean(y**2)
   10
   11 [dl_dw, dl_db] = tape.gradient(loss, [w, b])
   12 print(dl_dw)
   13 print(dl_db)

tf.Tensor(
[[0.07654426 0.67659944]
 [0.15308851 1.3531989 ]
 [0.22963277 2.0297983 ]], shape=(3, 2), dtype=float32)
tf.Tensor([0.07654426 0.67659944], shape=(2,), dtype=float32)
```

Graphs

- In the previous cases tf was run **eagerly**
- Graph execution benefits:
 - Portability outside python (mobile applications, embedded devices, ...)
 - Efficient execution
- Data structures that contain:
 - a set of **units of computation** (*tf.Operation*)
 - **units of data** that flow between operations (*tf.Tensor*)
- Graph can be saved, run, and restored all without the original Python code (and without Python interpreter)
- Graphs are also optimized, allowing the compiler to:
 - Separate sub-parts of a computation that are independent and split them between threads or devices
 - Simplify arithmetic operations by eliminating common subexpressions
 - [...]



Graphs

- In short, graphs let your tf **run fast, run in parallel, and run efficiently on multiple devices**
- We can still define our machine learning models in Python, and then **automatically construct graphs when we need them**
- *tf.function* take a canonical function as input and returns *Function*, a Python callable that build tf graphs from the python function
- On the outside, a *Function* looks like a regular function, but it encapsulates several *tf.Graphs* behind one API
- *tf.function* **applies** to a function and **all other functions it calls**

```
def inner_function(x, y, b):  
    x = tf.matmul(x, y)  
    x = x + b  
    return x  
  
# Use the decorator to make 'outer_function' a 'Function'.  
@tf.function  
def outer_function(x):  
    y = tf.constant([[2.0], [3.0]])  
    b = tf.constant(4.0)  
  
    return inner_function(x, y, b)  
  
# Note that the callable will create a graph that  
# includes 'inner_function' as well as 'outer_function'.  
outer_function(tf.constant([[1.0, 2.0]])).numpy()
```

Modules, layers, and models

- To do ML in TensorFlow, you are likely to need to define, save, and restore a model, which abstractly is
 - A function that computes something on tensors (a forward pass)
 - Some variables that can be updated in response to training
- Most models are made of layers
- **Layers are functions with a known mathematical structure** that can be reused and have **trainable variables**
- We will use **Keras**, high-level implementations of layers and models, built on the same foundational class [tf.Module](#)

Keras



- ***tf.keras.layers.Layer*** is the base class of all Keras layers, and it inherits from *tf.Module*
- Also provides a full-featured model class called ***tf.keras.Model***
- In general, Layer class is used to define inner computation blocks, and the Model class to define the outer model
- The Model class has the same API as Layer but exposing also:
 - Built-in training, evaluation, and prediction loops
 - The list of its inner layers, via the `model.layers` property.
 - Saving and serialization APIs
- See [Making new Layers and Models via subclassing](#)



Assignment

ResNet




ResNet - Intuition

- Experimentally, with the network **depth increasing**, **accuracy** gets saturated (which might be unsurprising) and then **degrades rapidly**
- Unintuitively more layers = less accuracy
- This is due to two reasons:
 - The vanishing of the gradient
 - The degradation problem



ResNet – Gradient vanishing

- Traditional activation functions (e.g., the hyperbolic tangent) have gradients in $(0,1]$
- **Backpropagation computes gradients by the chain rule**
- **Multiplying** n of these **small numbers** will cause the gradient to decrease exponentially
- As the gradient keeps flowing backward to the initial layers, this value keeps getting multiplied by each local gradient
- The vanishing of the gradient prevents the learning of the early layers




Local gradients of a neuron computed by backprop

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial x}$$
$$\frac{\partial L}{\partial y} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial y}$$


ResNet – Gradient vanishing

- A node in a computational graph can compute two things without even being aware of the rest of the graph :
 - the output of the node
 - the local gradient of the node.
- Local gradients of a node are the derivatives of the output of the node with respect to each of the inputs
- Essence of backpropagation: for a **single node**, we find the derivative of the output w.r.t a variable by **multiplying its local gradient with the upstream gradient** that we receive from the upstream step



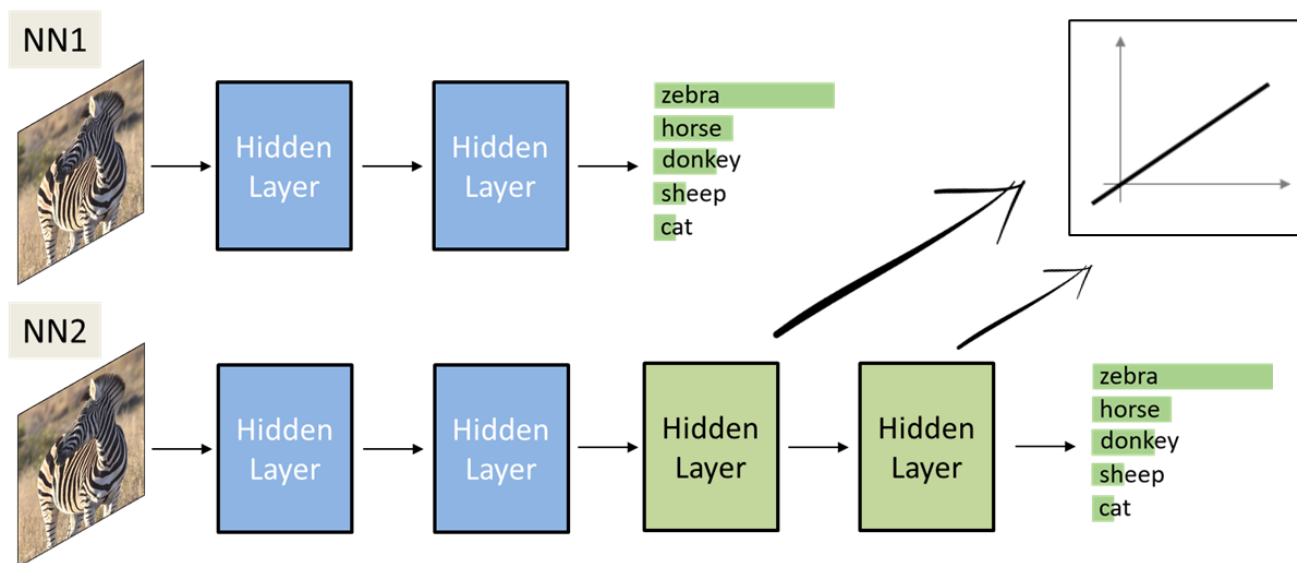
Chain rule

$$L(O_2(O_1(w_1))) \rightarrow \frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial O_2} \frac{\partial O_2}{\partial O_1} \frac{\partial O_1}{\partial w_1}$$




ResNet – Degradation problem

- **Theoretically**, the more layers you add to a neural network, the performance could either go up or stay the same, it should never go down
- Suppose NN1 in its *best possible state* (loss in global minimum)
- If we add more hidden layers (NN2), the new layers should learn the **identity function** $g(x) = x$ to preserve the current best state of the net
- Learning the identity function from scratch is extremely difficult: the huge number of weights and bias values must be modified to correspond to identity function



ResNet – Residual learning

- We introduce the *ResNet block*, a layer with **skip** connections

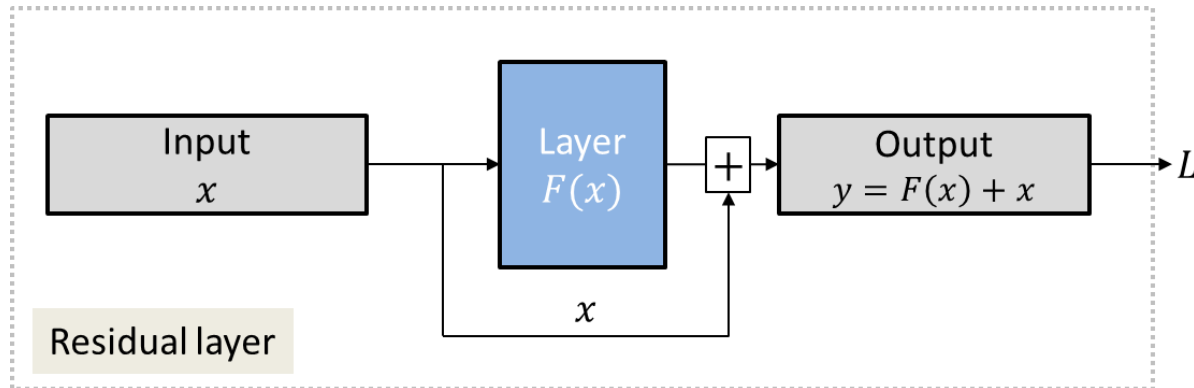


- Defining $y = F(x) + x$, the layer must learn $F(x) = y - x$, which is the **residual** mapping
- In this scenario, learning **the identity function** means learning the **zero** function $y = 0 + x = x$
- Easy to learn, set all weights to zero
- No vanishing gradient: $\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial L}{\partial y} (F'(x) + 1) = \frac{\partial L}{\partial y} + \frac{\partial L}{\partial y} F'(x)$


The **error signal (incoming gradient)** is passed through, unmodified by the **local gradient**

ResNet – Residual learning


- The architecture depends on the dimension of the output w.r.t. the input




- If the input dimension has the same dimension as the output the block implements an identity function
- Otherwise in the skip connection we insert a **1x1 convolutional block**



ResNet - 1x1 convolution




5x5x32



1x1x32

A 1x1 convolution will take an **element-wise product** and then apply a Relu non-linearity

ResNet - 1x1 convolution




Can be see a neuron having as input 32 number multiplied by weights \rightarrow Network In Network [1]

[1] Lin, Min, Qiang Chen, and Shuicheng Yan. "Network In Network." arXiv preprint arXiv:1312.4400 (2013).

ResNet - 1x1 convolution

1x1x32




5x5x32

3	5			

5x5

ResNet - 1x1 convolution

1x1x32




5x5x32

3	5	2		

5x5

ResNet - 1x1 convolution

1x1x32




5x5x32

3	5	2	7	


5x5

ResNet - 1x1 convolution




Using N different 1×1 convolutional filters we obtain volumes of different depth

ResNet - 1x1 convolution



Using N different 1x1 convolutional filters
we obtain volumes of different depth

ResNet - 1x1 convolution



Concerning ResNet, since there's no pooling layer within the residual block, the dimension is downscaled by 1×1 convolution with strides 2